



Machine Learning and Optimization Methods - IT3071

Assignment 1 - Practical test

IT NUMBER: IT22295538

NAME: I.L.M. BALASOORIYA

BATCH: DS.WE.0201

1.introduction

In order to address a multi-class classification problem, I created and built a deep learning model in this project utilizing TensorFlow's Keras API. The objective was to precisely categorize the data using 60 features into three categories: 'L', 'M', and 'H'. The model's performance was optimized by experimenting with different architectures and hyperparameters.

2. Data Preparation

I started by preprocessing the 480 rows and 17 columns that made up the dataset:

I started by looking over the data frame to see if there were any columns that were missing or had null values.

Encoding categorical variables: For categorical columns with more than two unique values per column, I employed one-hot encoding. used the label encoding, which had two distinct values, after that. Because of the classification, I specifically utilized label encoding for the Class column.

Columns dropping, I let go of the dummy One Way to Prevent Multicollinearity

Numerical feature scaling: To make sure that features had the same scale, MinMax scaling was used.

20% of the dataset was set aside for testing and the remaining 80% for training.

IT22295538.ipynb

colab.research.google.com/drive/1ISoGw6wdYAFoPY4I0DFHbLJPk387yZ#scrollTo=KWWhjrO78Cz1s

File Edit View Insert Runtime Tools Help

Files

- sample_data
- academic_dataset.csv

```
[21] data2.sample(5)

edhands  VisITedResources  AnnouncementsView  Discussion  ParentAnsweringSurvey  ParentschoolSatisfaction  StudentAbsenceDays  ...  Topic_Chemistry
28        60             19         50             1             1             0  ...  False
70        92             50         7             1             1             0  ...  False
72        80             58         66             0             0             0  ...  False
80        80             51         59             1             1             0  ...  False
72        51             42         24             1             0             1  ...  False

[22] data2.shape
(480, 67)

[24] for col in data2 :
      print(f'{col}\t\t{data2[col].unique()}') #use this loop to check, is there any non numeric values are still in the dataframe
StudentAbsenceDays      [0 1]
Class                   [0 1 2]
Nationality_Egypt        [False True]
Nationality_Iran          [False True]
Nationality_Iraq          [False True]
Nationality_Jordan        [False True]
Nationality_KW            [ True False]
Nationality_Lybia         [False True]
Nationality_Morocco        [False True]
Nationality_Palestine     [False True]
```

74.87 GB available

https://accounts.google.com/SignOutOptions?hl=en&continue=https://colab.research.google.com/&ec=GBRAqQM ing (1m 17s) <cell line: 12> > error_handler() > fit()

IT22295538.ipynb

colab.research.google.com/drive/1ISoGw6wdYAFoPY4I0DFHbLJPk387yZ#scrollTo=KWWhjrO78Cz1s

File Edit View Insert Runtime Tools Help

Files

- sample_data
- academic_dataset.csv

```
data.isnull().sum() #check whether there are any null values

gender      0
Nationality  0
PlaceofBirth  0
StageID     0
GradeID     0
SectionID   0
Topic       0
Semester    0
Relation    0
raisedhands  0
VisITedResources  0
AnnouncementsView  0
Discussion     0
ParentAnsweringSurvey  0
ParentschoolSatisfaction  0
StudentAbsenceDays  0
```

74.87 GB available

1m 34s completed at 6:34 PM

IT22295538.ipynb

colab.research.google.com/drive/1ISoGw6wdYAfoPY4I0DFHbLIPfk387yZ#scrollTo=KWWhjO78Cz1s

File Edit View Insert Runtime Tools Help

Files

- sample_data
- academic_dataset.csv

Code

```
[6] ParentschoolSatisfaction 0
     StudentAbsenceDays 0
     Class 0
dtype: int64

for col in data:
    print(f'{col}\t\t{data[col].unique()}') # check and print unique values per column

gender ['M' 'F']
Nationality ['KW' 'lebanon' 'Egypt' 'SaudiArabia' 'USA' 'Jordan' 'venezuela' 'Iran'
            'Tunis' 'Morocco' 'Syria' 'Palestine' 'Iraq' 'Lybia']
PlaceofBirth ['Kuwait' 'lebanon' 'Egypt' 'SaudiArabia' 'USA' 'Jordan' 'venezuela' 'Iran'
              'Tunis' 'Morocco' 'Syria' 'Iraq' 'Palestine' 'Lybia']
StageID ['lowerlevel' 'Middleschool' 'Highschool']
GradeID ['G-04' 'G-07' 'G-08' 'G-06' 'G-05' 'G-09' 'G-12' 'G-11' 'G-10' 'G-02']
SectionID ['A' 'B' 'C']
Topic ['IT' 'Math' 'Arabic' 'Science' 'English' 'Quran' 'Spanish' 'French'
       'History' 'Biology' 'Chemistry' 'Geology']
Semester ['F' 'S']
Relation ['Father' 'Mum']
raisedhands [15 20 10 30 40 42 35 50 12 70 19 5 62 36 55 69 60 2
            0 8 25 75 4 45 14 33 7 13 29 39 49 16 28 27 21 80
            17 65 22 11 1 3 100 6 90 77 24 66 23 82 72 51 85 87
            95 81 53 92 83 67 96 57 73 9 32 52 59 61 79 18 74 97
            41 71 98 78 89 88 86 76 99 84]
VisiTedResources [16 20 7 25 50 30 12 10 21 80 88 6 1 14 70 40 13 15 60 0 2 19 85 90
                 5 22 11 54 35 33 4 39 75 69 3 8 89 44 92 26 27 29 98 9 42 65 79 55
                 63 91 51 58 68 82 72 52 62 71 66 43 95 31 41 81 61 83 84 17 94 48 86 74
                 76 97 87 99 34 64 28 38 36 24 59 57 77 18 93 96 78]
AnnouncementsView [2 3 0 5 12 13 15 16 25 30 19 44 22 20 35 36 40 33 4 52 50 10 9 8]
```

1m 34s completed at 6:34PM

IT22295538.ipynb

colab.research.google.com/drive/1ISoGw6wdYAfoPY4I0DFHbLIPfk387yZ#scrollTo=KWWhjO78Cz1s

File Edit View Insert Runtime Tools Help All changes saved

Files

- sample_data
- academic_dataset.csv

Code

```
79 21 31 28 38 48 97 98 63 72 82 71 45 68 92 58 57 62]
ParentAnsweringSurvey ['Yes' 'No']
ParentschoolSatisfaction ['Good' 'Bad']
StudentAbsenceDays ['Under-7' 'Above-7']
Class ['M' 'L' 'H']

[12] #Label encoding

data.gender.replace({'F':0, 'M':1}, inplace=True)
data.Semester.replace({'F':0, 'S':1}, inplace=True)
data.Relation.replace({'Father':0, 'Mum':1}, inplace=True)
data.ParentAnsweringSurvey.replace({'No':0, 'Yes':1}, inplace=True)
data.ParentsSchoolSatisfaction.replace({'Bad':0, 'Good':1}, inplace=True)
data.StudentAbsenceDays.replace({'Under-7':0, 'Above-7':1}, inplace=True)
data.Class.replace({'M':0, 'L':1, 'H':2}, inplace=True)

#from sklearn.preprocessing import LabelEncoder
#le = LabelEncoder()
#data['Class'] = le.fit_transform(data['Class'])
#data['Semester'] = le.fit_transform(data['Semester'])
#data['gender'] = le.fit_transform(data['gender'])
#data['Relation'] = le.fit_transform(data['Relation'])
#data['ParentAnsweringSurvey'] = le.fit_transform(data['ParentAnsweringSurvey'])
#data['ParentsSchoolSatisfaction'] = le.fit_transform(data['ParentsSchoolSatisfaction'])
#data['StudentAbsenceDays'] = le.fit_transform(data['StudentAbsenceDays'])

[13] data.sample(5)
```

gender Nationality PlaceofBirth StageID GradeID SectionID Topic Semester Relation raisedhands VisiTedResources AnnouncementsView

1m 34s completed at 6:34PM

colab.research.google.com/drive/1ISoGhw6wdYAfoPY4I0DFHbLJPfk387yZ#scrollTo=KWbjrO78Cz1s

IT22295538.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Files

- sample_data
- academic_dataset.csv

Code

```
[13] 86 1 SaudiArabia SaudiArabia lowerlevel G-02 B IT 0 0 70 12
```

```
[20] #one-hot encoding
data2 = pd.get_dummies(data, columns=['Nationality', 'PlaceofBirth', 'StageID', 'GradeID', 'SectionID', 'Topic'])
```

```
data2.sample(5)
```

edhands	VisITedResources	AnnouncementsView	Discussion	ParentAnsweringSurvey	ParentschoolSatisfaction	StudentAbsenceDays	Topic_Chemistry
28	60	19	50	1	1	0	False
70	92	50	7	1	1	0	False
72	80	58	66	0	0	0	False
80	80	51	59	1	1	0	False
72	51	42	24	1	0	1	False

```
[22] data2.shape
```

```
(480, 67)
```

```
[24] for col in data2 :
      print(f'{col}\t\t{data2[col].unique()}') #use this loop to check, is there any non numeric values are still in the dataframe
```

1m 34s completed at 6:34 PM

colab.research.google.com/drive/1ISoGhw6wdYAfoPY4I0DFHbLJPfk387yZ#scrollTo=KWbjrO78Cz1s

IT22295538.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Files

- sample_data
- academic_dataset.csv

Code

```
[24] Topic_History [False True]
      Topic_IT [True False]
      Topic_Math [False True]
      Topic_Quran [False True]
      Topic_Science [False True]
      Topic_Spanish [False True]
```

```
[25] #Drop One Dummy Variable to Avoid Multicollinearity
data2.drop(['Nationality_Egypt', 'PlaceofBirth_Iran', 'StageID_MiddleSchool', 'GradeID_G-08', 'SectionID_A', 'Topic_History'], axis='columns', inplace=True)
```

```
data2.sample(5)
```

gender	Semester	Relation	raisedhands	VisITedResources	AnnouncementsView	Discussion	ParentAnsweringSurvey	ParentschoolSatisfaction
79	0	0	1	80	90	49	55	1
358	0	0	1	72	98	52	15	1
126	0	0	0	2	9	7	55	1
432	1	0	0	95	87	62	81	0
338	0	0	0	78	98	10	11	0

5 rows x 61 columns

```
[27] data2.shape
```

```
(480, 61)
```

1m 34s completed at 6:34 PM

IT22295538.ipynb - Colab

colab.research.google.com/drive/1ISoGw6wdYAfoPY4l0DFHbLJPfk387yZ#scrollTo=KWhjrO78Cz1s

IT22295538.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Comment Share Settings

RAM Disk Gemini

Files

- sample_data
- academic_dataset.csv

Code

```
[27] data2.shape
(480, 61)

[29] scaler = MinMaxScaler()

[30] col_to_scale = ['raisedhands', 'VisITedResources', 'AnnouncementsView', 'Discussion']

#Min-max scaling transforms the data to a fixed range
data2[col_to_scale] = scaler.fit_transform(data2[col_to_scale])

[32] data2.sample(5)
```

	gender	Semester	Relation	raisedhands	VisITedResources	AnnouncementsView	Discussion	ParentAnsweringSurvey	ParentschoolSatisfaction	S
416	1	0	0	0.98	0.909091	0.877551	0.714286	1	1	
432	1	0	0	0.95	0.878788	0.632653	0.816327	0	0	
260	1	1	0	0.10	0.171717	0.122449	0.132653	0	0	
30	0	0	0	0.35	0.808081	0.510204	0.704082	1	1	
233	0	1	1	0.32	0.808081	0.591837	0.459184	1	1	

5 rows x 61 columns

1m 34s completed at 6:34 PM

IT22295538.ipynb - Colab

colab.research.google.com/drive/1ISoGw6wdYAfoPY4l0DFHbLJPfk387yZ#scrollTo=KWhjrO78Cz1s

IT22295538.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Comment Share Settings

RAM Disk Gemini

Files

- sample_data
- academic_dataset.csv

Code

```
[326] #Split the dataset into training and validation sets
from sklearn.model_selection import train_test_split

[327] x = data2.drop('Class',axis='columns')
y = data2.Class

x.sample(5)
```

	gender	Semester	Relation	raisedhands	VisITedResources	AnnouncementsView	Discussion	ParentAnsweringSurvey	ParentschoolSatisfaction	S
330	1	0	0	0.40	0.070707	0.510204	0.408163	0	1	
150	1	1	0	0.80	0.808081	0.520408	0.591837	1	1	
191	1	1	0	0.15	0.252525	0.377551	0.122449	1	1	
302	0	0	1	0.11	0.202020	0.214286	0.224490	0	0	
362	1	0	0	0.90	0.989899	0.418367	0.377551	1	1	

5 rows x 60 columns

Double-click (or enter) to edit

```
[329] x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.2,random_state=47)

[330] x_train.shape
```

1m 34s completed at 6:34 PM

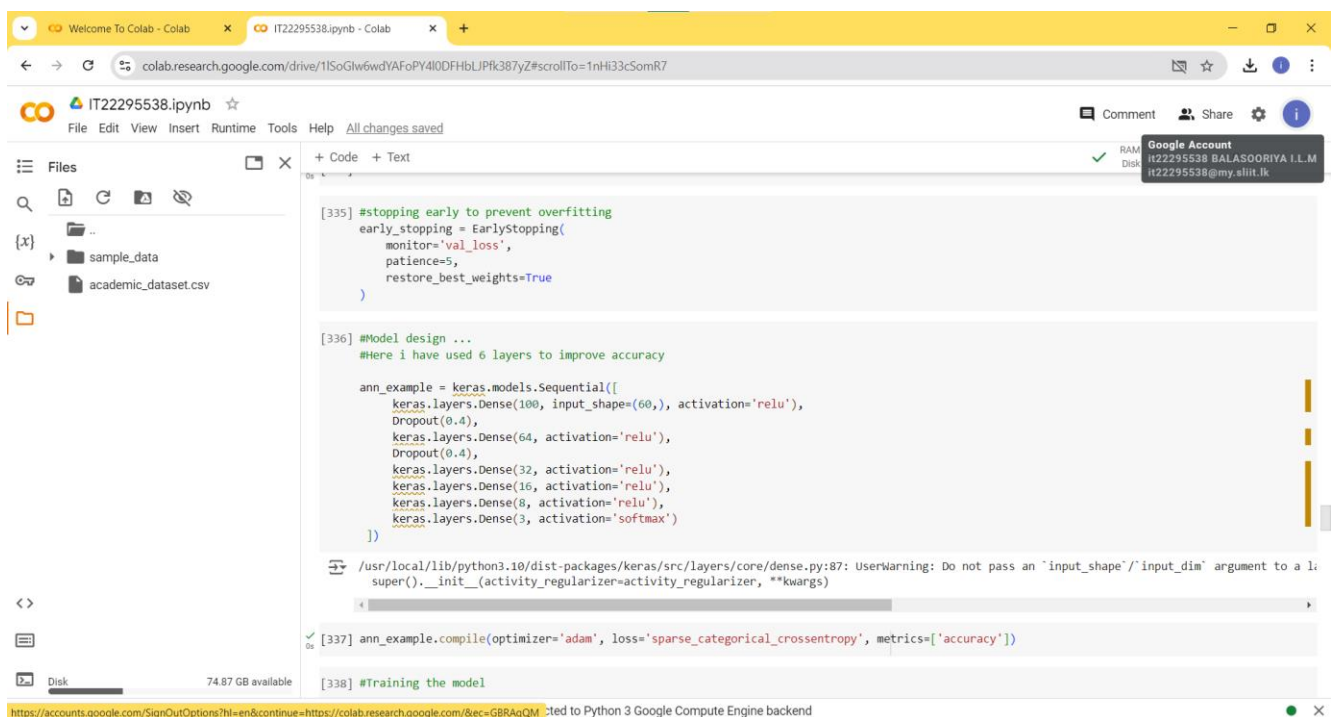
3. Model designing

I experimented with different model architectures. The final architecture consisted of:

Dense Layers: 5 layers with 'relu' activations and a 'Softmax' output layer for classification into 3 classes.

Dropout Layers: Dropout was introduced to prevent overfitting, with a rate of 0.4.

Optimizer: Adam was selected as the optimizer, and sparse categorical cross-entropy was used as the loss function and accuracy as metrics.



The screenshot shows a Google Colab notebook interface. The left sidebar displays the file explorer with a folder named 'sample_data' and a file named 'academic_dataset.csv'. The main code area contains the following Python code:

```
[335] #stopping early to prevent overfitting
early_stopping = EarlyStopping(
    monitor='val_loss',
    patience=5,
    restore_best_weights=True
)

[336] #Model design ...
#Here i have used 6 layers to improve accuracy

ann_example = keras.models.Sequential([
    keras.layers.Dense(100, input_shape=(60,), activation='relu'),
    Dropout(0.4),
    keras.layers.Dense(64, activation='relu'),
    Dropout(0.4),
    keras.layers.Dense(32, activation='relu'),
    keras.layers.Dense(16, activation='relu'),
    keras.layers.Dense(8, activation='relu'),
    keras.layers.Dense(3, activation='softmax')
])

/usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer that does not support it.
super().__init__(activity_regularizer=activity_regularizer, **kwargs)

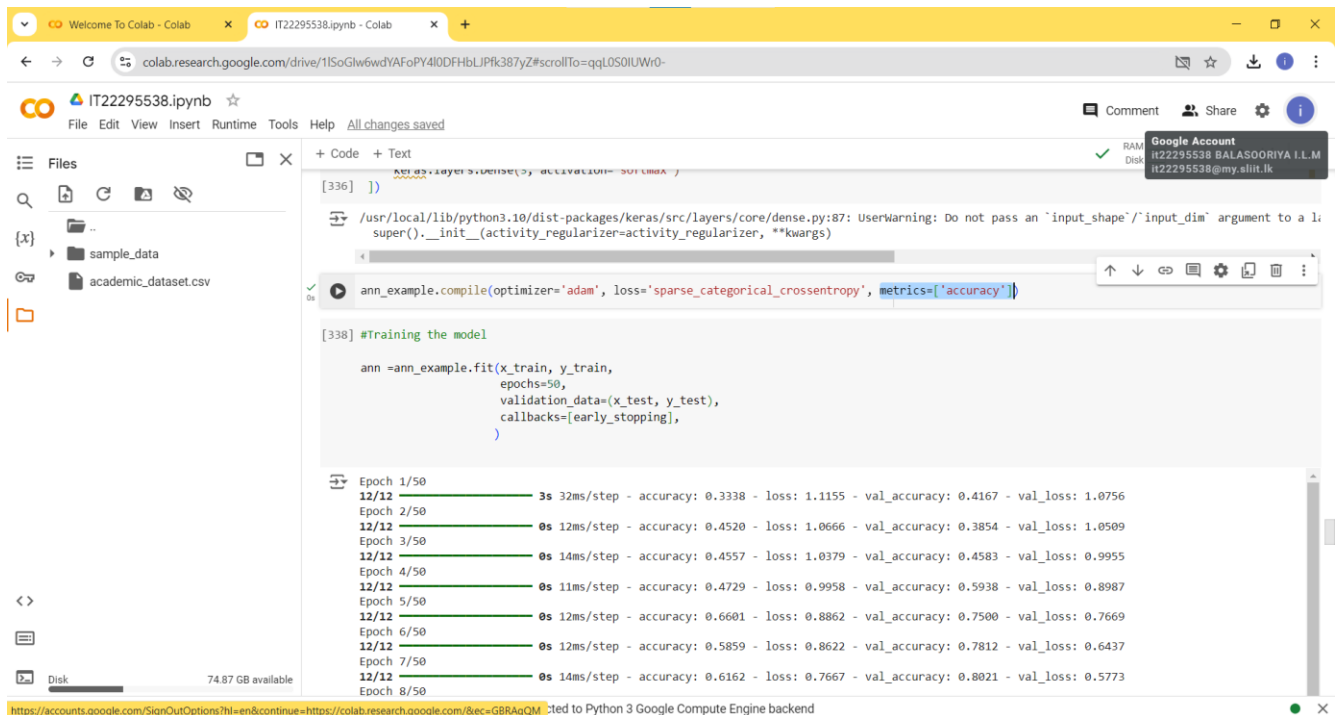
[337] ann_example.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

[338] #Training the model
```

The bottom of the notebook shows the status bar with '74.87 GB available' and a link to the Google Account page.

4. Training

Training loss and accuracy were logged during the training process, which was monitored for 50 epochs with early stopping to avoid overfitting.



```
[336] ]  
/usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer.  
super().__init__(activity_regularizer=activity_regularizer, **kwargs)  
ann_example.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])  
[338] #Training the model  
ann = ann_example.fit(x_train, y_train,  
                      epochs=50,  
                      validation_data=(x_test, y_test),  
                      callbacks=[early_stopping],  
                      )  
Epoch 1/50  
12/12 3s 32ms/step - accuracy: 0.3338 - loss: 1.1155 - val_accuracy: 0.4167 - val_loss: 1.0756  
Epoch 2/50  
12/12 0s 12ms/step - accuracy: 0.4520 - loss: 1.0666 - val_accuracy: 0.3854 - val_loss: 1.0509  
Epoch 3/50  
12/12 0s 14ms/step - accuracy: 0.4557 - loss: 1.0379 - val_accuracy: 0.4583 - val_loss: 0.9955  
Epoch 4/50  
12/12 0s 11ms/step - accuracy: 0.4729 - loss: 0.9958 - val_accuracy: 0.5938 - val_loss: 0.8987  
Epoch 5/50  
12/12 0s 12ms/step - accuracy: 0.6601 - loss: 0.8862 - val_accuracy: 0.7500 - val_loss: 0.7669  
Epoch 6/50  
12/12 0s 12ms/step - accuracy: 0.5859 - loss: 0.8622 - val_accuracy: 0.7812 - val_loss: 0.6437  
Epoch 7/50  
12/12 0s 14ms/step - accuracy: 0.6162 - loss: 0.7667 - val_accuracy: 0.8021 - val_loss: 0.5773  
Epoch 8/50
```

5. Evaluation

The trained model was evaluated on the test set, and the following metrics were calculated:

Test Loss: 0.3940761983394623

Test Accuracy: 0.8541666865348816

Precision, Recall, F1-score: These metrics were calculated for each class using `classification_report()`.

colab.research.google.com/drive/1ISoGhw6wdYAFoPY4I0DFHbLJPik387yZ#scrollTo=ZCizUjyCtobY

IT22295538.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Comment Share

Google Account
IT22295538 BALASOORIYA I.L.M
it22295538@my.slit.it

Files

Code + Text

```
12/12 0s 6ms/step - accuracy: 0.8124 - loss: 0.4010 - val_accuracy: 0.8123 - val_loss: 0.4023
Epoch 30/50
12/12 0s 6ms/step - accuracy: 0.8277 - loss: 0.3898 - val_accuracy: 0.8542 - val_loss: 0.4020
Epoch 31/50
12/12 0s 7ms/step - accuracy: 0.8067 - loss: 0.4640 - val_accuracy: 0.8542 - val_loss: 0.3951

#evaluating the model

test_loss, test_accuracy = ann_example.evaluate(x_test, y_test)
print(f'Test Loss: {test_loss}')
print(f'Test Accuracy: {test_accuracy}')

3/3 0s 5ms/step - accuracy: 0.8763 - loss: 0.3614
Test Loss: 0.3940761983394623
Test Accuracy: 0.8541666865348816

[ ] predicted = ann_example.predict(x_test)

3/3 0s 4ms/step

[ ] predicted_classes = [np.argmax(i) for i in predicted]

[ ] print(classification_report(y_test, predicted_classes))
```

colab.research.google.com/drive/1ISoGhw6wdYAFoPY4I0DFHbLJPik387yZ#scrollTo=ZCizUjyCtobY

IT22295538.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Comment Share

Google Account
IT22295538 BALASOORIYA I.L.M
it22295538@my.slit.it

Files

Code + Text

```
3/3 0s 5ms/step - accuracy: 0.8763 - loss: 0.3614
Test Loss: 0.3940761983394623
Test Accuracy: 0.8541666865348816

[ ] predicted = ann_example.predict(x_test)

3/3 0s 4ms/step

[ ] predicted_classes = [np.argmax(i) for i in predicted]

print(classification_report(y_test, predicted_classes))

precision    recall  f1-score   support

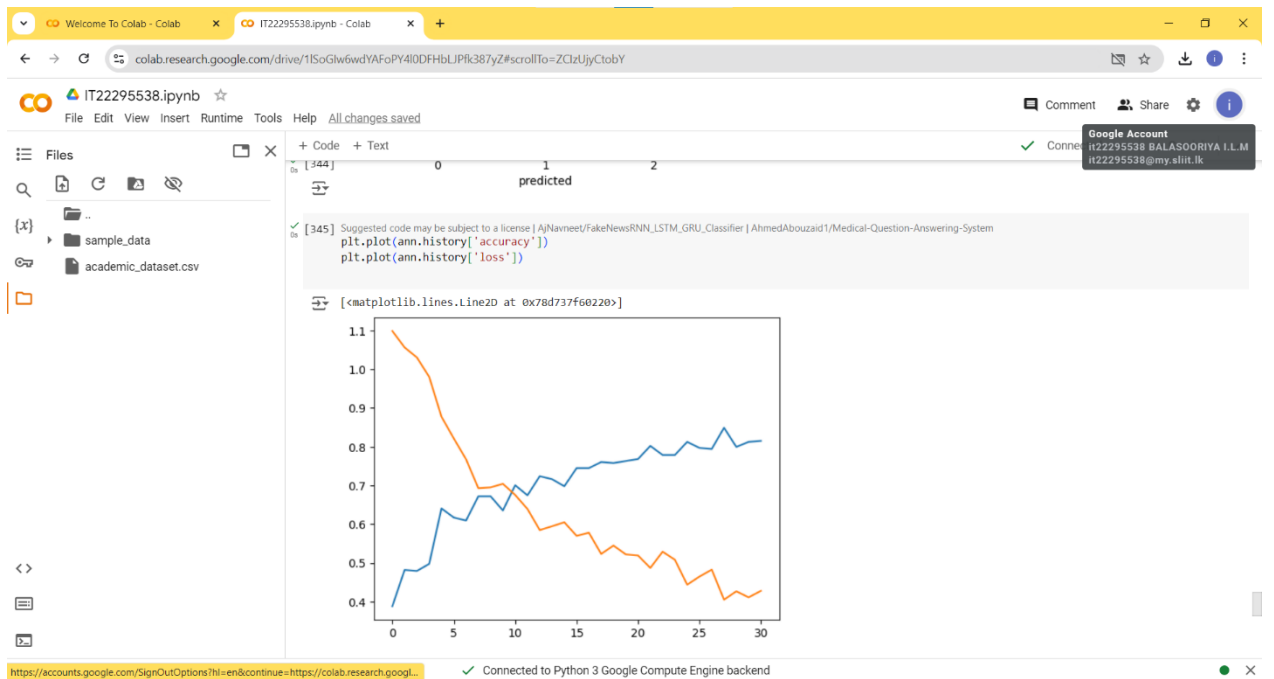
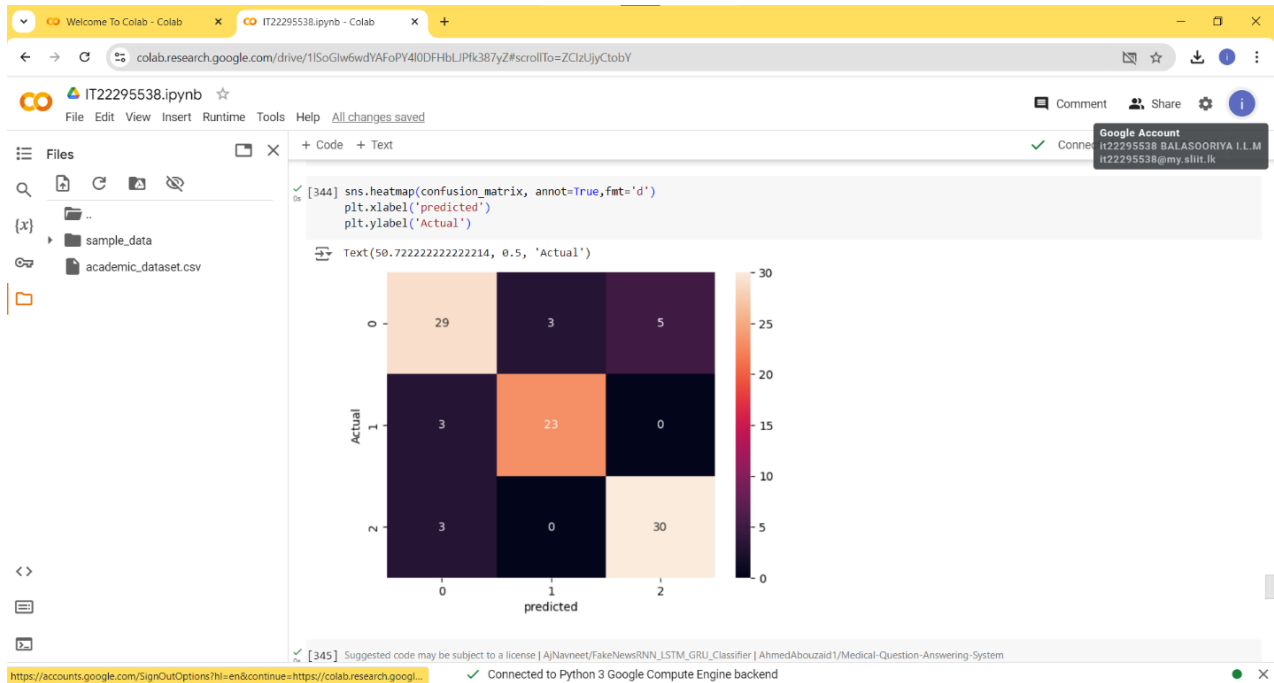
0     0.83     0.78     0.81     37
1     0.88     0.88     0.88     26
2     0.86     0.91     0.88     33

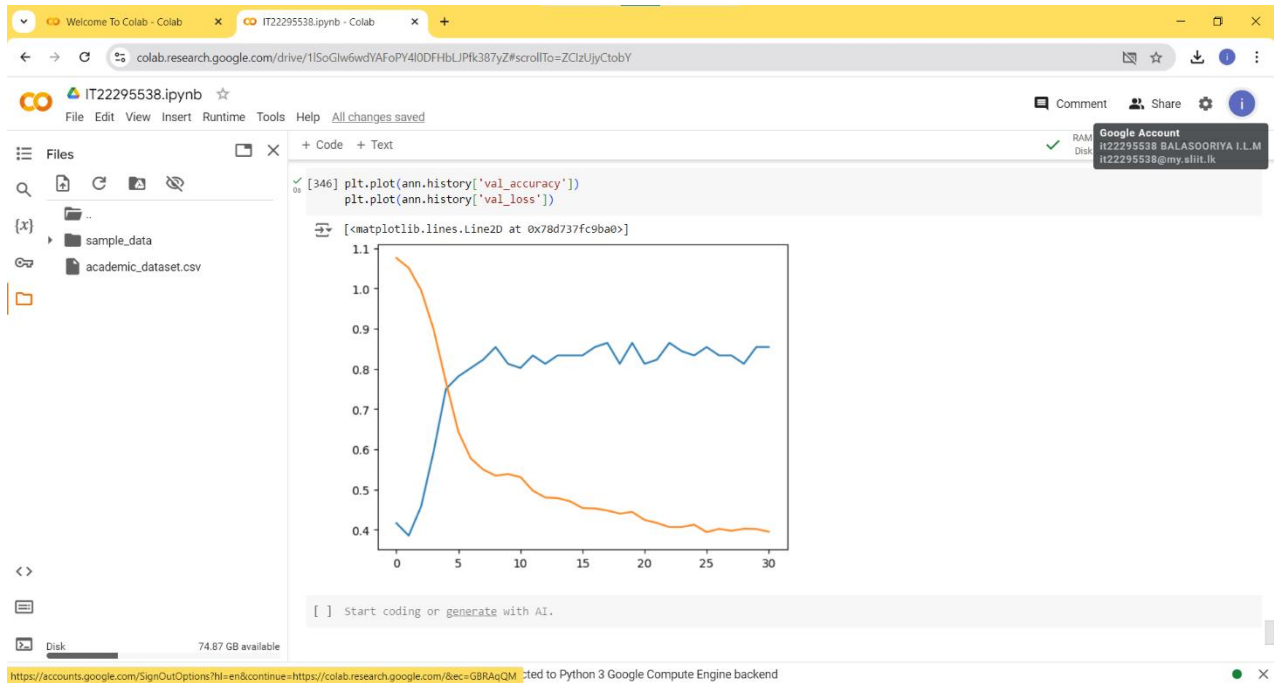
accuracy          0.86
macro avg          0.86
weighted avg       0.85

[ ] confusion_matrix = confusion_matrix(y_test, predicted_classes)

[ ] sns.heatmap(confusion_matrix, annot=True,fmt='d')
plt.xlabel('predicted')
plt.ylabel('Actual')
```

Initializing Python 3 Google Compute Engine backend





6. Hyperparameter Tuning

To further optimize performance, I experimented with different hyperparameters such as:

Batch size: 16, 32, 64

Learning rate: 0.001, 0.01, 0.1

The best results were achieved with a batch size of 16 and a learning rate of 0.001.

colab.research.google.com/drive/1ISoGw6wdYAFoPY4I0DFHbLIPk387yZ#scrollTo=KWhjO78Cz1s

IT22295538.ipynb

```
param_grid = {
    'batch_size': [16, 32, 64],
    'learning_rate': [0.001, 0.01, 0.1],
}

grid = ParameterGrid(param_grid)

best_accuracy = 0
best_params = {}

for params in grid:
    print(f"Testing parameters: {params}")

    # Model with current parameters
    ann_example = keras.models.Sequential([
        keras.layers.Dense(100, input_shape=(60,)), activation='relu',
        Dropout(0.4),
        keras.layers.Dense(64, activation='relu'),
        Dropout(0.4),
        keras.layers.Dense(32, activation='relu'),
        keras.layers.Dense(16, activation='relu'),
        keras.layers.Dense(8, activation='relu'),
        keras.layers.Dense(3, activation='softmax')
    ])

    optimizer = tf.keras.optimizers.Adam(learning_rate=params['learning_rate'])
    ann_example.compile(optimizer=optimizer,
                        loss='sparse_categorical_crossentropy',
                        metrics=['accuracy'])

    # Train the model
    ann = ann_example.fit(x_train, y_train,
                          batch_size=params['batch_size'],
                          epochs=50,
                          validation_data=(x_test, y_test),
                          verbose=0)

    # Evaluate the model
    evaluation = ann.evaluate(x_test, y_test, verbose=0)
    validation_accuracy = evaluation[1]

    # Update best parameters if current model is better
    if validation_accuracy > best_accuracy:
        best_accuracy = validation_accuracy
        best_params = params
        print(f"*****Results*****")
        print(f"Best parameters: {best_params}")
        print(f"Best validation accuracy: {best_accuracy}")

    print(f"*****Results*****")
    print(f"Best parameters: {best_params}")
    print(f"Best validation accuracy: {best_accuracy}")
```

Executing (1m 1s) <cell line: 12> -> error_handler() -> fit() -> error_handler() -> evaluate() -> error_handler() -> _call__() -> _call() -> call_function() -> _call_flat() -> call_flat() -> call_function() -> quick_execute()

colab.research.google.com/drive/1ISoGw6wdYAFoPY4I0DFHbLIPk387yZ#scrollTo=KWhjO78Cz1s

IT22295538.ipynb

```
best_accuracy = validation_accuracy
best_params = params
print(f"*****Results*****")
print(f"Best parameters: {best_params}")
print(f"Best validation accuracy: {best_accuracy}")
```

Testing parameters: {'batch_size': 16, 'learning_rate': 0.001}

/usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an 'input_shape'/'input_dim' argument to a layer. When using Sequential models, prefer using a super().__init__(activity_regularizer=activity_regularizer, **kwargs)

Testing parameters: {'batch_size': 16, 'learning_rate': 0.01}

/usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an 'input_shape'/'input_dim' argument to a layer. When using Sequential models, prefer using a super().__init__(activity_regularizer=activity_regularizer, **kwargs)

Testing parameters: {'batch_size': 16, 'learning_rate': 0.1}

/usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an 'input_shape'/'input_dim' argument to a layer. When using Sequential models, prefer using a super().__init__(activity_regularizer=activity_regularizer, **kwargs)

Testing parameters: {'batch_size': 32, 'learning_rate': 0.001}

/usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an 'input_shape'/'input_dim' argument to a layer. When using Sequential models, prefer using a super().__init__(activity_regularizer=activity_regularizer, **kwargs)

Testing parameters: {'batch_size': 32, 'learning_rate': 0.01}

/usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an 'input_shape'/'input_dim' argument to a layer. When using Sequential models, prefer using a super().__init__(activity_regularizer=activity_regularizer, **kwargs)

Testing parameters: {'batch_size': 32, 'learning_rate': 0.1}

/usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an 'input_shape'/'input_dim' argument to a layer. When using Sequential models, prefer using a super().__init__(activity_regularizer=activity_regularizer, **kwargs)

Testing parameters: {'batch_size': 64, 'learning_rate': 0.001}

/usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an 'input_shape'/'input_dim' argument to a layer. When using Sequential models, prefer using a super().__init__(activity_regularizer=activity_regularizer, **kwargs)

Testing parameters: {'batch_size': 64, 'learning_rate': 0.01}

/usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an 'input_shape'/'input_dim' argument to a layer. When using Sequential models, prefer using a super().__init__(activity_regularizer=activity_regularizer, **kwargs)

Testing parameters: {'batch_size': 64, 'learning_rate': 0.1}

/usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an 'input_shape'/'input_dim' argument to a layer. When using Sequential models, prefer using a super().__init__(activity_regularizer=activity_regularizer, **kwargs)

*****Results*****

Best parameters: {'batch_size': 16, 'learning_rate': 0.001}

Best validation accuracy: 0.385416656725592

Start coding or generate with AI.

1m 41s completed at 7:02 PM

7. Findings

Overfitting was mitigated using dropout layers and early stopping, which resulted in smoother convergence of the model during training.

8. Challenges Faced

Choosing the right architecture: Experimenting with different architectures and hyperparameters required time to find an optimal solution.

Early stopping: Fine-tuning early stopping criteria took several trials to prevent underfitting while ensuring good generalization.

9. Lessons Learned

Dropout and early stopping are effective techniques to combat overfitting.

Hyperparameter tuning can significantly impact model performance, and automating this process using grid search or random search can save time.

Monitoring both loss and accuracy metrics is crucial to assess model performance over time.